
Project Milestone

Abhishek Goswami
Microsoft
Redmond, WA 98052
agoswami@microsoft.com

Abstract

This document details the progress made towards the project milestone of Course CS224n, Winter 2019.

1 Introduction

Machine comprehension (MC) and question answering (QA) tasks have gained significant interest in the past few years, with several end-to-end models showing promising results. A key factor in recent advancements has been the use of neural attention mechanisms, which extract useful signal by exploiting the notion of *matching*.

Several attention approaches have been proposed in literature. Chen et al [3] propose a uni-directional attention mechanism whereby the query attends to the context paragraph. In BiDAF, Seo et al [10] introduce bi-directional attention flow to obtain query-aware context representations. Wang et al [12] note that question-aware passage representations have limited knowledge of the context itself, and motivate self-matching attention to directly match the question-aware passage representation against itself.

Another key tenet of the proposed techniques is to use a model to process sequential inputs. This is typically done in the form of an embedding encoder layer. While recurrent neural networks have been the model of choice for this, recent work by Yu et al [13] propose using a combination of convolution and self-attention mechanisms instead.

In this paper we explore two novel extensions. First, we observe that given the input embeddings, most existing models dive straight into the encoding layer. Attention is an afterthought, that gets applied only in later layers of the model. A problem with such representations is that it strains the embedding encoder layer, with the rest of the modeling layers all stacked on top. We propose adding a ‘Embedding Attention Layer’ as a form of self-attention over the word and character input embeddings. Second, we propose having a contextual embed layer that uses a combination of recurrent layers and convolution layers, with the motivation of bringing the best of both worlds in the contextual embedding space. We show that adding these two extensions is indeed helpful, and show our comparisons with respect to the BiDAF model [10]

The remainder of the paper is organized as follows. In Section 2 we introduce our model. Section 3 presents the experimental results from our modeling techniques. In Section 4 we survey some related work in the area of machine comprehension. Finally, we present our conclusions in Section 5

2 Model

In this section, we first formulate the machine comprehension problem and then describe the model.

2.1 Problem Statement

The machine comprehension task considered in this paper is as follows. Given a context paragraph with T words, $C = \{c_1, c_2, \dots, c_T\}$ and a query sentence with J words, $Q = \{q_1, q_2, \dots, q_J\}$, output a span $S = \{c_i, c_{i+1}, \dots, c_{i+j}\}$ from the original paragraph C that satisfactorily answers the question. Section 3.3 describes two metrics that are widely used in literature for evaluating this task. We use d to represent the hidden size used by several layers of the model.

Table 1: An example of a machine comprehension task.

Question	Economy, Energy and Tourism is one of the what?
Context	Subject Committees are established at the beginning of each parliamentary session, and again the members on each committee reflect the balance of parties across Parliament. Typically each committee corresponds with one (or more) of the departments (or ministries) of the Scottish Government. The current Subject Committees in the fourth Session are: Economy, Energy and Tourism; Education and Culture; Health and Sport; Justice; Local Government and Regeneration; Rural Affairs, Climate Change and Environment; Welfare Reform; and Infrastructure and Capital Investment
Answer	current Subject Committees

2.2 Model Overview

Several state-of-the-art machine comprehension models have a similar structure. They have (a) an embedding layer (b) an embedding encoder layer (c) an attention flow layer (d) a model encoder layer and (e) an output layer.

We introduce two novel extensions to this structure. One, we add a Embedding Attention Layer between the input Embedding Layer and the Embedding Encoder Layer, with the goal of introducing early attention in the modeling process. Second, for the Embedding Encoder Layer we use a combination of recurrent and convolution operations to make it rich with both sequential and local interactions. Our machine comprehension model is thus a hierarchical multi-stage process consisting of six layers.

1. Embedding Layer.
2. Embedding Attention Layer.
3. Embedding Encoder Layer.
4. Attention Flow Layer.
5. Model Encoder Layer.
6. Output Layer.

The details of each of the layers are as follows.

1. Embedding Layer. In this layer we mix character embeddings with word embeddings.

For character embeddings, we use a method similar to that proposed by Kim et al [5]. We first convert a word to its character indices. We then pad (or truncate) each word so it has length m_{word} . For each of these characters we lookup a dense character embedding (which has shape e_{char}). To combine the character embeddings, we use 1-dimensional convolutions over m_{word} using e_{char} as the input channel size and e_{word} as the output channel size. The output of the CNN are max-pooled over the entire width to obtain a fixed-size vector of shape e_{word} for each word.

For word embeddings, we use pre-trained word vectors from GloVe [7] to obtain the fixed embedding for each word. The size of the word embeddings is e_{word} which is the same as the shape of the character-level embeddings for each word.

The concatenation of the character and word embeddings is passed to a Highway Network [11]. We do this for both the context sentence C and also the question Q . So now we have two matrices $C \in \mathbb{R}^{T,d}$ and $Q \in \mathbb{R}^{J,d}$ corresponding to the context and question respectively.

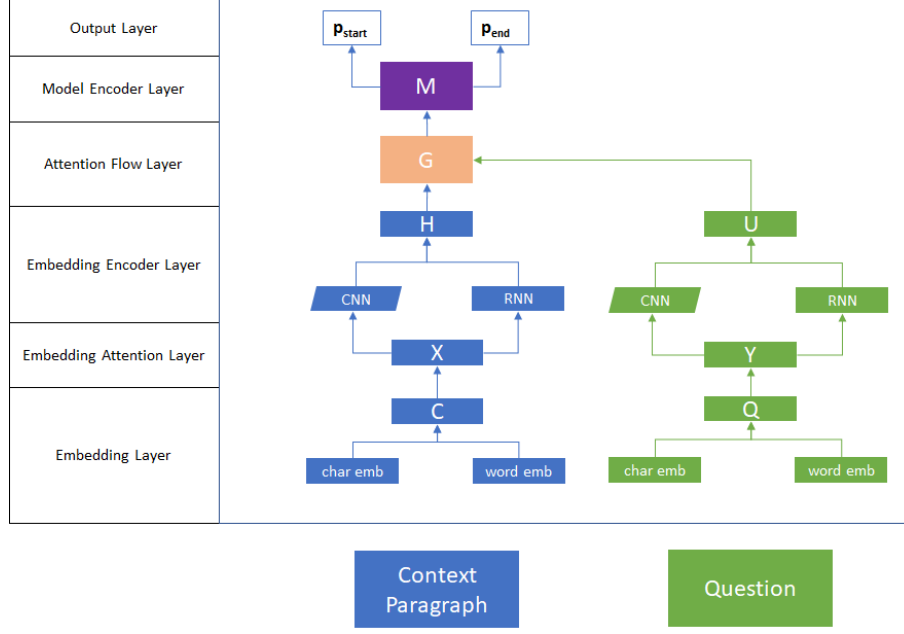


Figure 1: Model architecture

2. Embedding Attention Layer. The motivation for adding this layer is to attend to the embeddings provided by the previous layer. This layer starts early attention to the word and character embeddings. As before, we do this for both the context sentence C and also the question Q , and concatenate the results with the embedding layer output giving us two matrices $\mathbf{X} \in \mathbb{R}^{T, 2d}$ and $\mathbf{Y} \in \mathbb{R}^{J, 2d}$ corresponding to the context and question respectively.

3. Embedding Encoder Layer. The purpose of this layer is to encode the relationships between the embeddings provided by the previous layers. On one hand we want to model the temporal interactions between words. For this we use a bi-directional LSTM. This results in two matrices of shape $(T, 2d)$ and $(J, 2d)$ corresponding to the context and question respectively.

We also model local interactions between the embeddings output by the embedding attention layer. We use 1-dimensional convolutions over the sequence length using $2d$ as both the input and output channel size. We do this using a kernel size 1, which results in two matrices of shape $(T, 2d)$ and $(J, 2d)$ corresponding to the context and question respectively.

The concatenation of the RNN and CNN layers gives us two matrices $\mathbf{H} \in \mathbb{R}^{T, 4d}$ and $\mathbf{U} \in \mathbb{R}^{J, 4d}$ respectively.

4. Attention Flow Layer. We also add a bi-directional attentional flow layer introduced by Seo et al [10]. The main idea is that attention should flow both ways - from the context to the question and from the question to the context. The attention flow layer also fuses the information between the context and the query words.

The inputs to the layer are contextual vector representations of the context \mathbf{H} and the query \mathbf{U} . The outputs of the layer is $\mathbf{G} \in \mathbb{R}^{T, 16d}$ which is a query-aware vector representations of the context words, along with the embeddings from the previous layer.

5. Model Encoder Layer. This layer encodes the query-aware representations of the context words. The input is \mathbf{G} , and the output is matrix \mathbf{M} , which captures the interaction among the context words conditioned on the query. We use two layers of bi-directional LSTM, with hidden size d for each direction. Matrix $\mathbf{M} \in \mathbb{R}^{T, 2d}$ is then passed to the Output Layer.

6. Output Layer. This layer is application specific. For the QA task being explored in this project, we need to find a sub-phrase of the context to answer the query. The phrase is derived by predicting

the start and end indices of the phrase in the paragraph. The output layer produces two probability distribution $\mathbf{p}_{start}, \mathbf{p}_{end} \in \mathbb{R}^N$ corresponding to each position in the context.

$$\mathbf{p}_{start} = \text{softmax}(\mathbf{W}_{start}[\mathbf{G}; \mathbf{M}]). \quad (1)$$

$$\mathbf{p}_{end} = \text{softmax}(\mathbf{W}_{end}[\mathbf{G}; \mathbf{M}']). \quad (2)$$

where $\mathbf{M}' \in \mathbb{R}^{T, 2d}$ is a matrix obtained by applying a bi-directional LSTM to \mathbf{M} .

2.3 Model Training and Scoring

Training. We define the training loss as the sum of the negative log-likelihood (cross-entropy) loss for the start and end locations. So for a (context, question) pair with *start* index $i \in \{1, 2, \dots, T\}$ and *end* index $j \in \{1, 2, \dots, T\}$

$$\text{loss} = -\log \mathbf{p}_{start}(i) - \log \mathbf{p}_{end}(j). \quad (3)$$

During training, we average across the batch and use the Adadelta optimizer [14] to minimize the loss.

Scoring. At test time, we chose the pair (i,j) of indices that maximizes $\mathbf{p}_{start}(i) \cdot \mathbf{p}_{end}(j)$ subject to $i \leq j$ and $j - i + 1 \leq L_{max}$, where L_{max} is a hyperparameter which sets the maximum length of a predicted answer.

No Answer. We adopt the approach proposed by Levy et al [6]. We prepend a OOV token to the beginning of each sequence. The model outputs \mathbf{p}_{start} and \mathbf{p}_{end} soft-predictions as usual. When discretizing a prediction, if $\mathbf{p}_{start}(0) \cdot \mathbf{p}_{end}(0)$ is greater than any predicted answer span, the model predicts no-answer. Otherwise the model predicts the highest probability span. Note, this approach also allows us to predict a per-example confidence score that the question is unanswerable.

3 Experiment

In this section, we conduct experiments to study the performance of our models. We will benchmark our models on the Stanford Question Answering Dataset (SQuAD) 2.0 [9], considered to be one of the most competitive datasets in QA tasks. We also provide some implementation details for our models and present the main results.

3.1 Dataset

We consider the Stanford Question Answering Dataset (SQuAD) 2.0 [9] for machine comprehension. Our model is given a paragraph, and a question about that paragraph, as input. The goal is to answer the question correctly. There are around 150k questions, and roughly half of the questions cannot be answered using the provided paragraph.

3.1.1 Data splits

The official SQuAD dataset has three splits: train, dev and test. The train and dev sets are publicly available and the test set is entirely secret. For this project we use a custom dev and test set obtained by splitting the official dev set in half.

To summarize we have the following data splits:

- **Train.** 129,941 examples. All taken from the official SQuAD 2.0 training set.
- **Dev.** 6,078 examples. Roughly half of the official dev set, randomly selected.
- **Test.** 5,915 examples. The remaining examples from the official dev set.

From now on we refer to these splits as the train set, dev set and test set respectively. We will use the train set to train the model. We report the performance metrics on the dev set.

3.2 Training Details

The model architecture used for this task is shown in Figure 1.

For the Embedding Layer, m_{word} is set to 16. e_{char} and e_{word} are set to 64 and 300 respectively. We use one 1D filter for the CNN char embedding with a kernel size of 5. The hidden state d of the model is 100. For the convolutions in the Embedding Encoder Layer, we use a set of 4 stacked CNN layers, each with input/output channels as $2d$ with a kernel size of 1. We use dropout as a form of regularization across all the six layers in our model. Table 4 shows the effect of dropout on our model performance. We use the Adadelta optimizer [14] with a learning rate of 0.5 which is kept fixed. While training we use a batch size of 64. When scoring, L_{max} is set to 15.

We implement our model in Python using PyTorch [2]. The experiments are carried out on a Azure Data Science Virtual Machine (DSVM) [1] which has a NVIDIA Tesla K80 GPU.

3.3 Metric Details

We measure performance via two metrics: Exact Match (EM) and the F1 score.

- **Exact Match** is a binary measure (i.e. true/false) of whether the system output matches the ground truth answer exactly.
- **F1** is the harmonic mean of precision and recall.
- When a question has no answer, both the F1 and EM score are 1 if the model predicts no-answer, and 0 otherwise.
- For questions that do have answers, when evaluating on the dev or test sets, we take the maximum F1 and EM scores across the three human-provided answers for that question.

3.4 Results

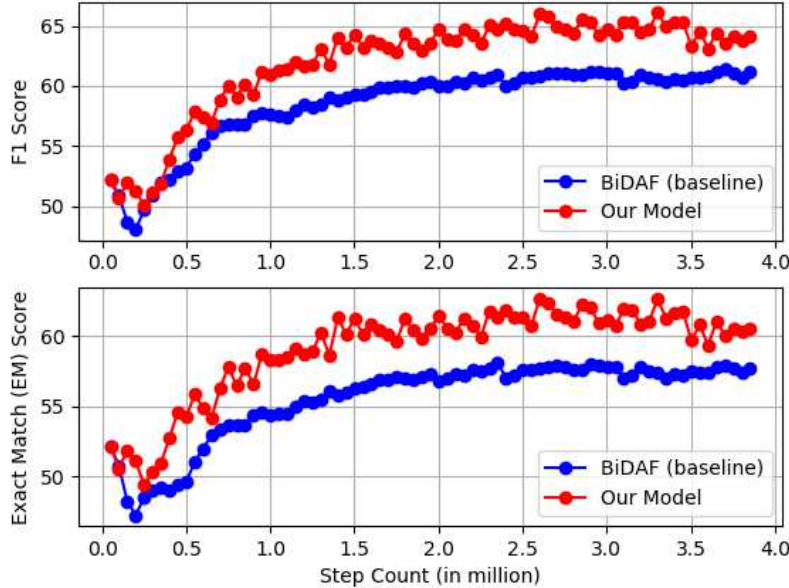


Figure 2: Comparison of F1 and EM scores

Table 5 shows the comparison between our model and the baseline. As per the original BiDAF model, we include a character-level embedding layer using character-level convnets. This gives us a very strong baseline to compare with.

Table 2: Comparing our model with the baseline

	EM	F1
BiDAF with character embedding (baseline)	59.47	62.46
Our Model	62.64	66.10

Table 3: Results from Ablation Study

	EM	F1
No Embedding Attention Layer	60.93	64.34
No CNN layers inside the Embedding Encoder Layer	59.64	63.10
No character embedding in the Embedding Layer	59.47	62.46
Freezing both the character and word embeddings in the Embedding Layer	62.19	65.39

3.5 Ablations

Table 3 shows the performance of the model and its ablations on the SQuAD dev set. Having an Embedding Attention Layer helps model performance. This validates our hypothesis that adding attention layers early in the model stack should help performance. For ablating the effect of the CNN layers, we experiment by removing the CNN layers from the Embedding Encoder Layer. CNN layers prove to be critical with a drop of 3 points on both metrics. As noted by Seo et al [10], having character embeddings in the Embedding Layer contributes towards model performance whereby word-level embeddings represent the semantics of each word as a whole, while char-level embeddings better handle out-of-vocab (OOV) or rare words. Interestingly we also see that freezing the char-level and word-level embeddings gives us slightly lower performance. It seems the model gives slightly better results if we allow the backpropagation to happen all the way through the embedding layer, and the fine-tuned embeddings generalize quite well to unseen data in the dev set.

Table 3 shows the effect of dropout rates. With low dropout rates, the model was overfitting. High drop-out rates help in preventing overfitting, but lead to lower EM/F1 scores. We settle on 0.2 as the dropout rate since it gave the best results.

4 Related Work

Machine comprehension (MC) and question answering (QA) tasks have gained significant interest in the past few years. Overall, the models and techniques that work best of these tasks fall into two categories.

One, techniques that leverage pre-trained contextual embeddings (PCE). Examples of such PCE-based techniques are ELMo [8] and BERT [4]. The core idea of such techniques is that in order to represent a piece of text, we should use word embeddings that depend on the context in which the word appears in the text. This is typically achieved by pretraining the weights on a large-scale language modeling dataset, and using the pre-trained weights for the initial model layers.

Secondly, there are the several end-to-end, non-PCE models which have shown promising results. Examples of such techniques are BiDAF [10], R-NET [12] and QANet [13].

Table 4: Effect of dropout

	EM	F1
No Dropout	60.41	63.46
Dropout = 0.1	62.06	65.39
Dropout = 0.2 (chosen)	62.64	66.10
Dropout = 0.3	59.84	63.81
Dropout = 0.4	61.10	64.09

Table 5: Preliminary Results

	Dev Set		Test Set	
	EM	F1	EM	F1
master	59.25	62.28	59.47	62.46
baselinemodel_selfsimilaritybeforeattention	58.83	62.05	57.47	60.87
baselinemodel_selfsimilarityafterattention	59.45	62.95	59.45	62.95
conflict_without_for_loops	55.67	59.19	55.67	59.19
embedding_with_self_similarity	59.64	63.10	59.64	63.10
embedding_with_self_difference	52.19	52.19	52.19	52.19
conv_rnn_embed_layer	60.93	64.34	60.93	64.34
embedding_with_self_similarity_conv_rnn (0.0)	62.19	65.26	60.41	63.46
embedding_with_self_similarity_conv_rnn (0.1)	62.19	65.26	62.06	65.39
embedding_with_self_similarity_conv_rnn (0.2)	62.19	65.26	62.19	65.26
embedding_with_self_similarity_conv_rnn (0.3)	62.19	65.26	59.84	63.18
embedding_with_self_similarity_conv_rnn (0.4)	62.19	65.26	61.10	64.09
embedding_with_self_similarity_conv_rnn (both unfrozen)	62.19	65.26	62.64	66.10

5 Conclusion

So far we have seen that character-level word embeddings help improve the baseline model. As next step we plan to implement the Base Attentional Model (Section ??) and the RNN-Conv Contextual Embedding Model (Section ??)

References

- [1] Azure data science virtual machines. <https://azure.microsoft.com/en-us/services/virtual-machines/data-science-virtual-machines/>.
- [2] Pytorch. <https://pytorch.org/>.
- [3] D. Chen, J. Bolton, and C. D. Manning. A thorough examination of the CNN/Daily Mail reading comprehension task. In *Association for Computational Linguistics (ACL)*, 2016.
- [4] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- [5] Y. Kim, Y. Jernite, D. Sontag, and A. M. Rush. Character-aware neural language models. In *Thirtieth AAAI Conference on Artificial Intelligence*, 2016.
- [6] O. Levy, M. Seo, E. Choi, and L. Zettlemoyer. Zero-shot relation extraction via reading comprehension. *arXiv preprint arXiv:1706.04115*, 2017.
- [7] J. Pennington, R. Socher, and C. Manning. Glove: Global vectors for word representation. In *Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP)*, pages 1532–1543, 2014.
- [8] M. E. Peters, M. Neumann, M. Iyyer, M. Gardner, C. Clark, K. Lee, and L. Zettlemoyer. Deep contextualized word representations. *arXiv preprint arXiv:1802.05365*, 2018.
- [9] P. Rajpurkar, R. Jia, and P. Liang. Know what you don’t know: Unanswerable questions for squad. *arXiv preprint arXiv:1806.03822*, 2018.
- [10] M. Seo, A. Kembhavi, A. Farhadi, and H. Hajishirzi. Bidirectional attention flow for machine comprehension. *arXiv preprint arXiv:1611.01603*, 2016.
- [11] R. K. Srivastava, K. Greff, and J. Schmidhuber. Highway networks. *arXiv preprint arXiv:1505.00387*, 2015.

- [12] W. Wang, N. Yang, F. Wei, B. Chang, and M. Zhou. Gated self-matching networks for reading comprehension and question answering. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, volume 1, pages 189–198, 2017.
- [13] A. W. Yu, D. Dohan, M.-T. Luong, R. Zhao, K. Chen, M. Norouzi, and Q. V. Le. Qanet: Combining local convolution with global self-attention for reading comprehension. *arXiv preprint arXiv:1804.09541*, 2018.
- [14] M. D. Zeiler. Adadelta: an adaptive learning rate method. *arXiv preprint arXiv:1212.5701*, 2012.